

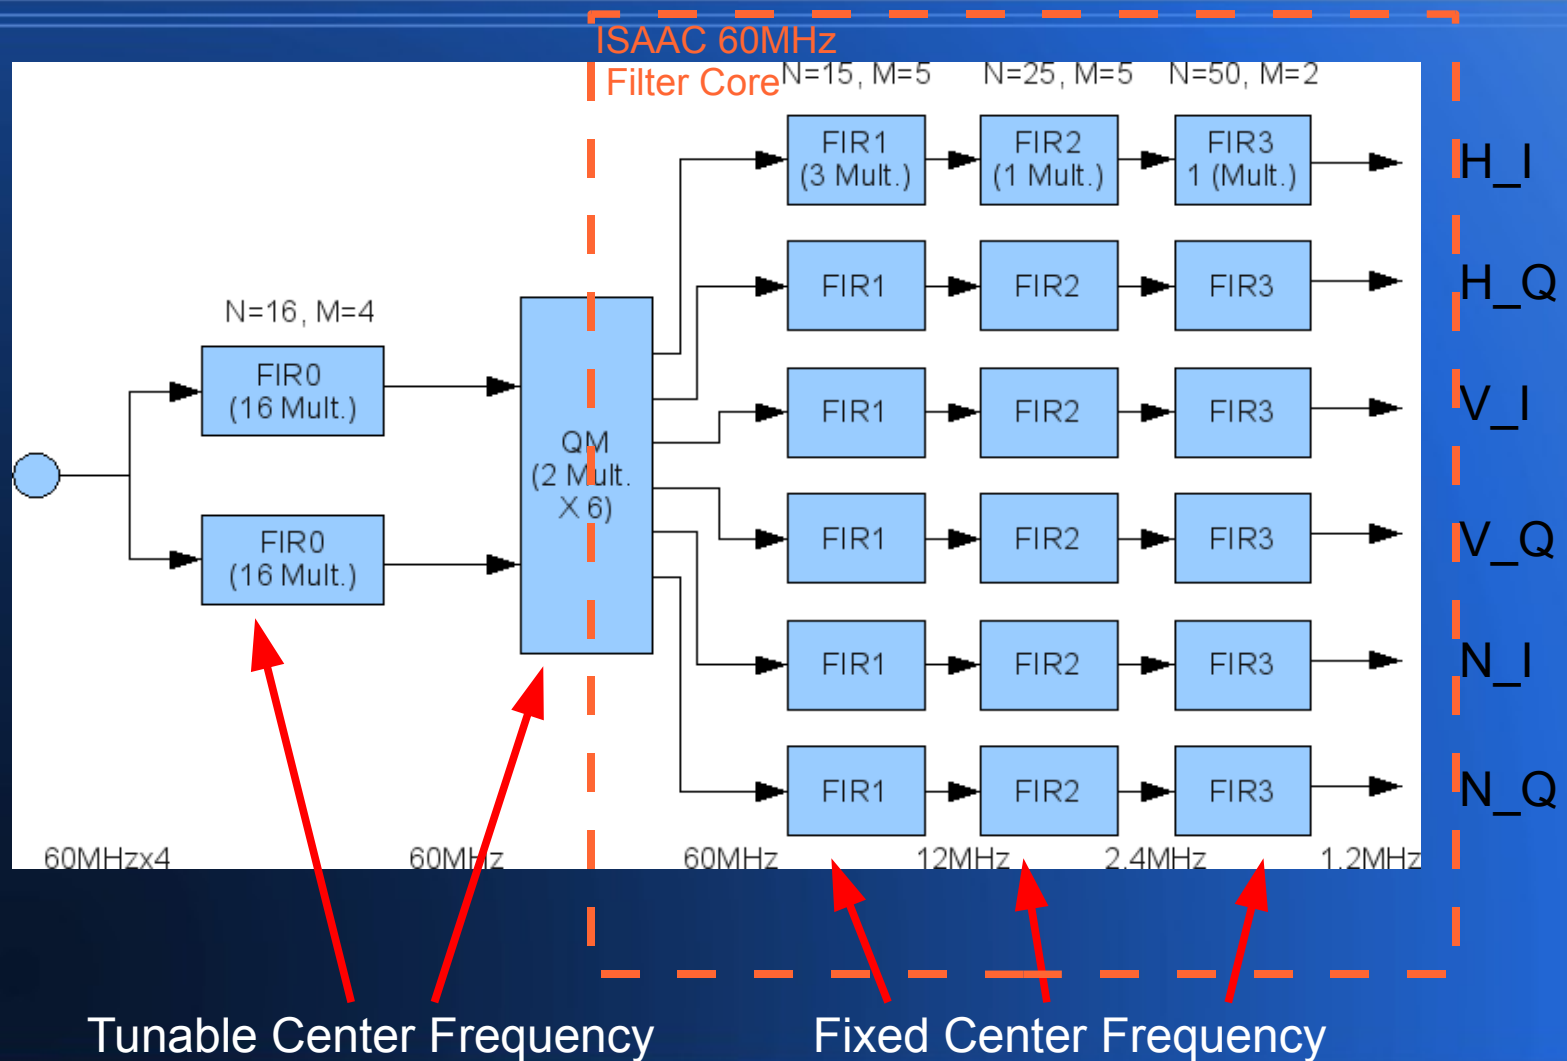
240MHz Digital Filter V2-3000 Implementation

Jason Zheng, Kayla Nguyen,
Charles Le, Yutao He


Outline

- Overview
- FIRDECIM implementation
- Coefficient Bank Implementation
- QM implementation
- FIR and QM Tuning
- Verification Results
- Synthesis Results

Overview



FIRDECIM



- FIR followed by a decimation stage
- Commonly referred to as polyphase FIR or firdecim
- N = number of taps, M = decimation factor

Implementation Choices

- Plain FIR + Decimation
- Simple Polyphase Filter
- Serial FIR (distributed arithmetic) + Decimation
- ISAAC Thread-Oriented FIR
- Considerations:
 - Speed and Size
 - Ease of Implementation, Modification, and Verification
 - Center Frequency Tuning
 - SEU mitigation (coefficient correctness)

Plain FIR + Decimation


- Assume no input buffer
- multiplier clock to input rate ratio of R
- N taps
- Need $\text{ceil}(N/R)$ multipliers
- Most calculations are thrown away due to decimation



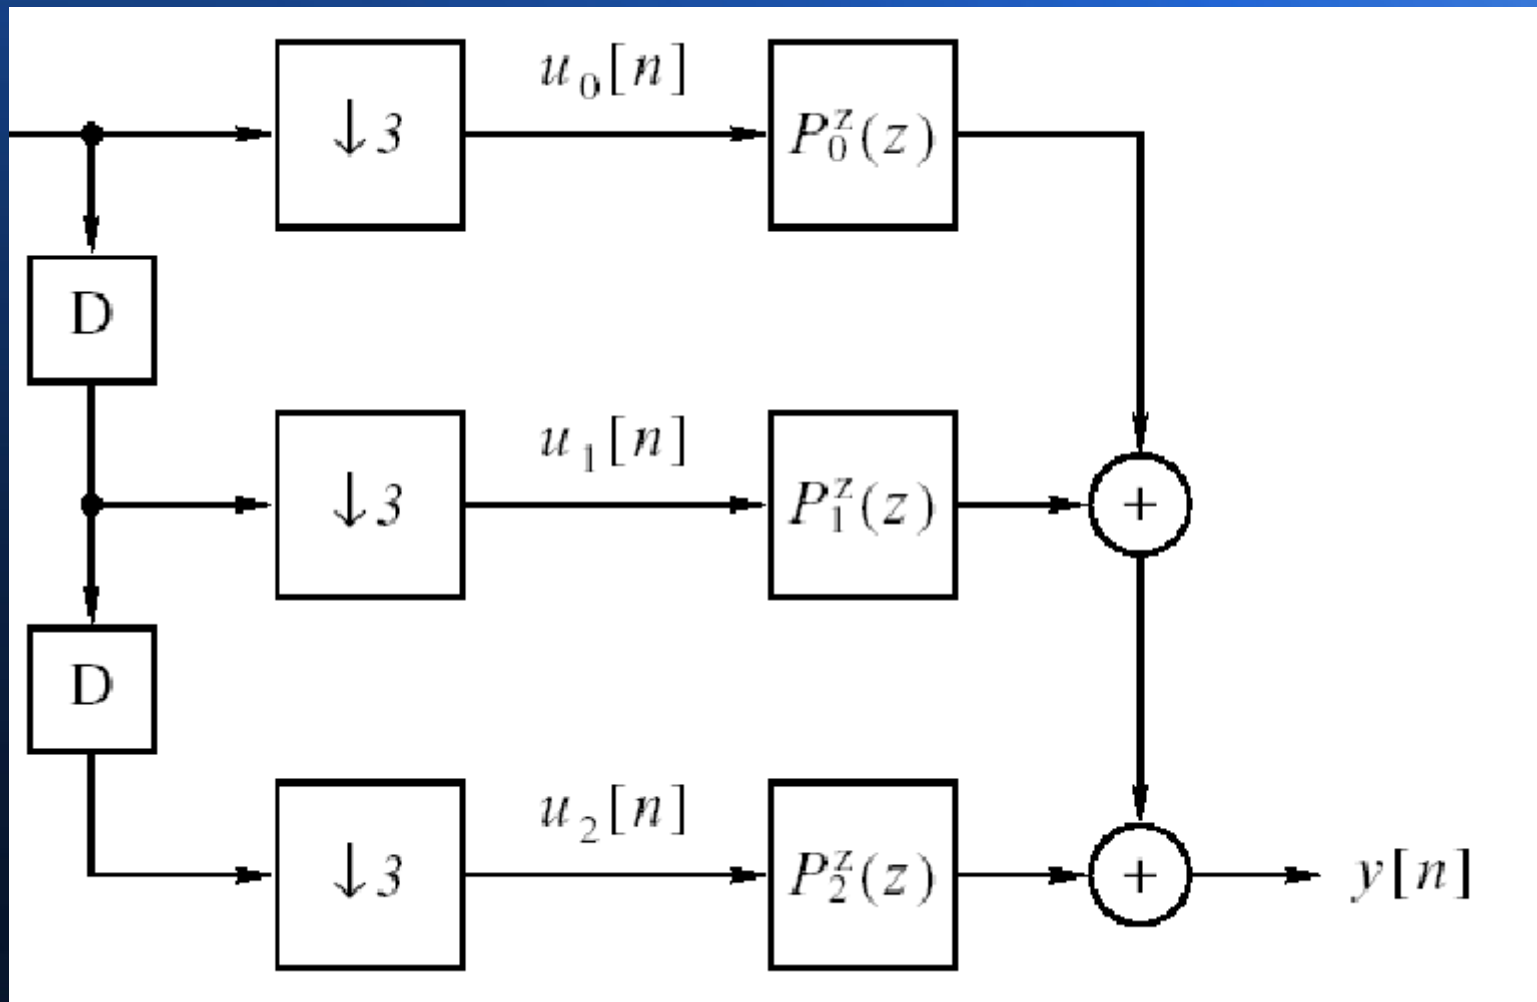
Simple Polyphase Filter

- Decimate first, then filter
- Break N -tap filter into M sub-filters
- No wasted calculations
- Require $\max(N/M, M)$ multipliers, at least 1 multiplier per sub-filter

Polyphase Filter – Switched Input View



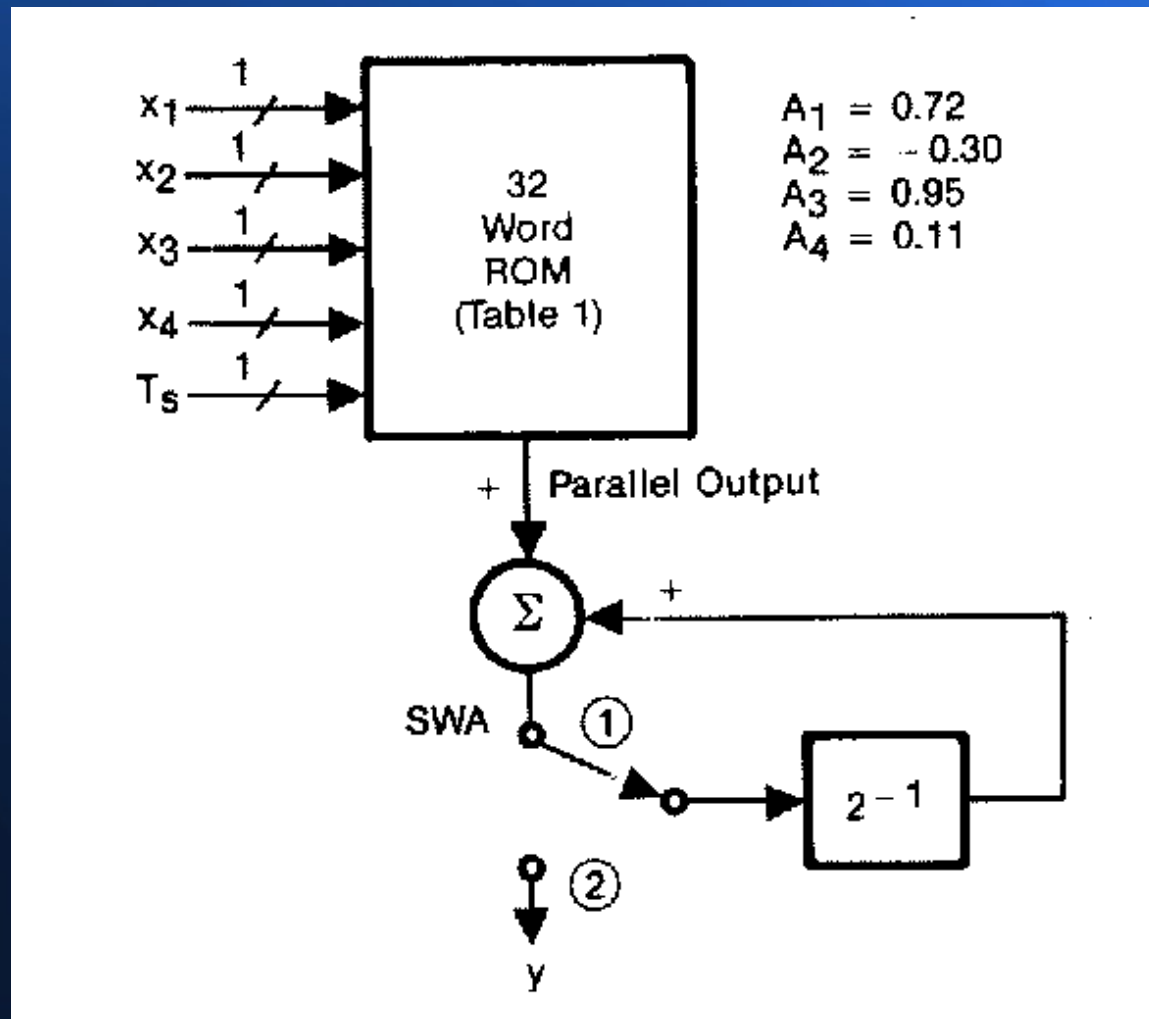
Polyphase Filter - Delayed Inputs View



Serial FIR

- Calculate N multiplications all at once, but bit-wise.
- Latency depends on fixed point width K
- Multipliers are replaced by LUTs (2^N depth)
- Internal Clock rate = $K * \text{Input rate}$
- Need input buffer for higher input rates
- Tuning = rewrite all LUTs

Serial FIR (From White's Paper)



Our Approach – Thread Oriented

- Based on textbook polyphase filter
- Generalize multiplier sharing with decimation factor M and clock ratio R
- Multiplier needed is $\text{ceil}(N / (M * R))$
- R is typically very high in a multi-stage firdecim design, esp in later stages.
- Suits well for “bottom-heavy” filters with single clock domain

Our Approach – Thread View

Each output is computed as an individual thread with a MAC unit

	x0	x1	x2	x3	x4	x5	x6	x7	x8	x9	x10	x11	x12	x13	x14	x15	x16	x17	x18	x19	x20	x21	x22	x23	x24	x25	x26
output1	*e15	*e14	*e13	*e12	*e11	*e10	*e9	*e8	*e7	*e6	*e5	*e4	*e3	*e2	*e1	*e0											
		*e15	*e14	*e13	*e12	*e11	*e10	*e9	*e8	*e7	*e6	*e5	*e4	*e3	*e2	*e1	*e0										
			*e15	*e14	*e13	*e12	*e11	*e10	*e9	*e8	*e7	*e6	*e5	*e4	*e3	*e2	*e1	*e0									
				*e15	*e14	*e13	*e12	*e11	*e10	*e9	*e8	*e7	*e6	*e5	*e4	*e3	*e2	*e1	*e0								
output2					*e15	*e14	*e13	*e12	*e11	*e10	*e9	*e8	*e7	*e6	*e5	*e4	*e3	*e2	*e1	*e0							
						*e15	*e14	*e13	*e12	*e11	*e10	*e9	*e8	*e7	*e6	*e5	*e4	*e3	*e2	*e1	*e0						
							*e15	*e14	*e13	*e12	*e11	*e10	*e9	*e8	*e7	*e6	*e5	*e4	*e3	*e2	*e1	*e0					
								*e15	*e14	*e13	*e12	*e11	*e10	*e9	*e8	*e7	*e6	*e5	*e4	*e3	*e2	*e1	*e0				
output3									*e15	*e14	*e13	*e12	*e11	*e10	*e9	*e8	*e7	*e6	*e5	*e4	*e3	*e2	*e1	*e0			
										*e15	*e14	*e13	*e12	*e11	*e10	*e9	*e8	*e7	*e6	*e5	*e4	*e3	*e2	*e1	*e0		
											*e15	*e14	*e13	*e12	*e11	*e10	*e9	*e8	*e7	*e6	*e5	*e4	*e3	*e2	*e1	*e0	
												*e15	*e14	*e13	*e12	*e11	*e10	*e9	*e8	*e7	*e6	*e5	*e4	*e3	*e2	*e1	*e0
output4													*e15	*e14	*e13	*e12	*e11	*e10	*e9	*e8	*e7	*e6	*e5	*e4	*e3	*e2	*e1
														*e15	*e14	*e13	*e12	*e11	*e10	*e9	*e8	*e7	*e6	*e5	*e4	*e3	*e2
															*e15	*e14	*e13	*e12	*e11	*e10	*e9	*e8	*e7	*e6	*e5	*e4	*e3
output4																*e15	*e14	*e13	*e12	*e11	*e10	*e9	*e8	*e7	*e6	*e5	*e4

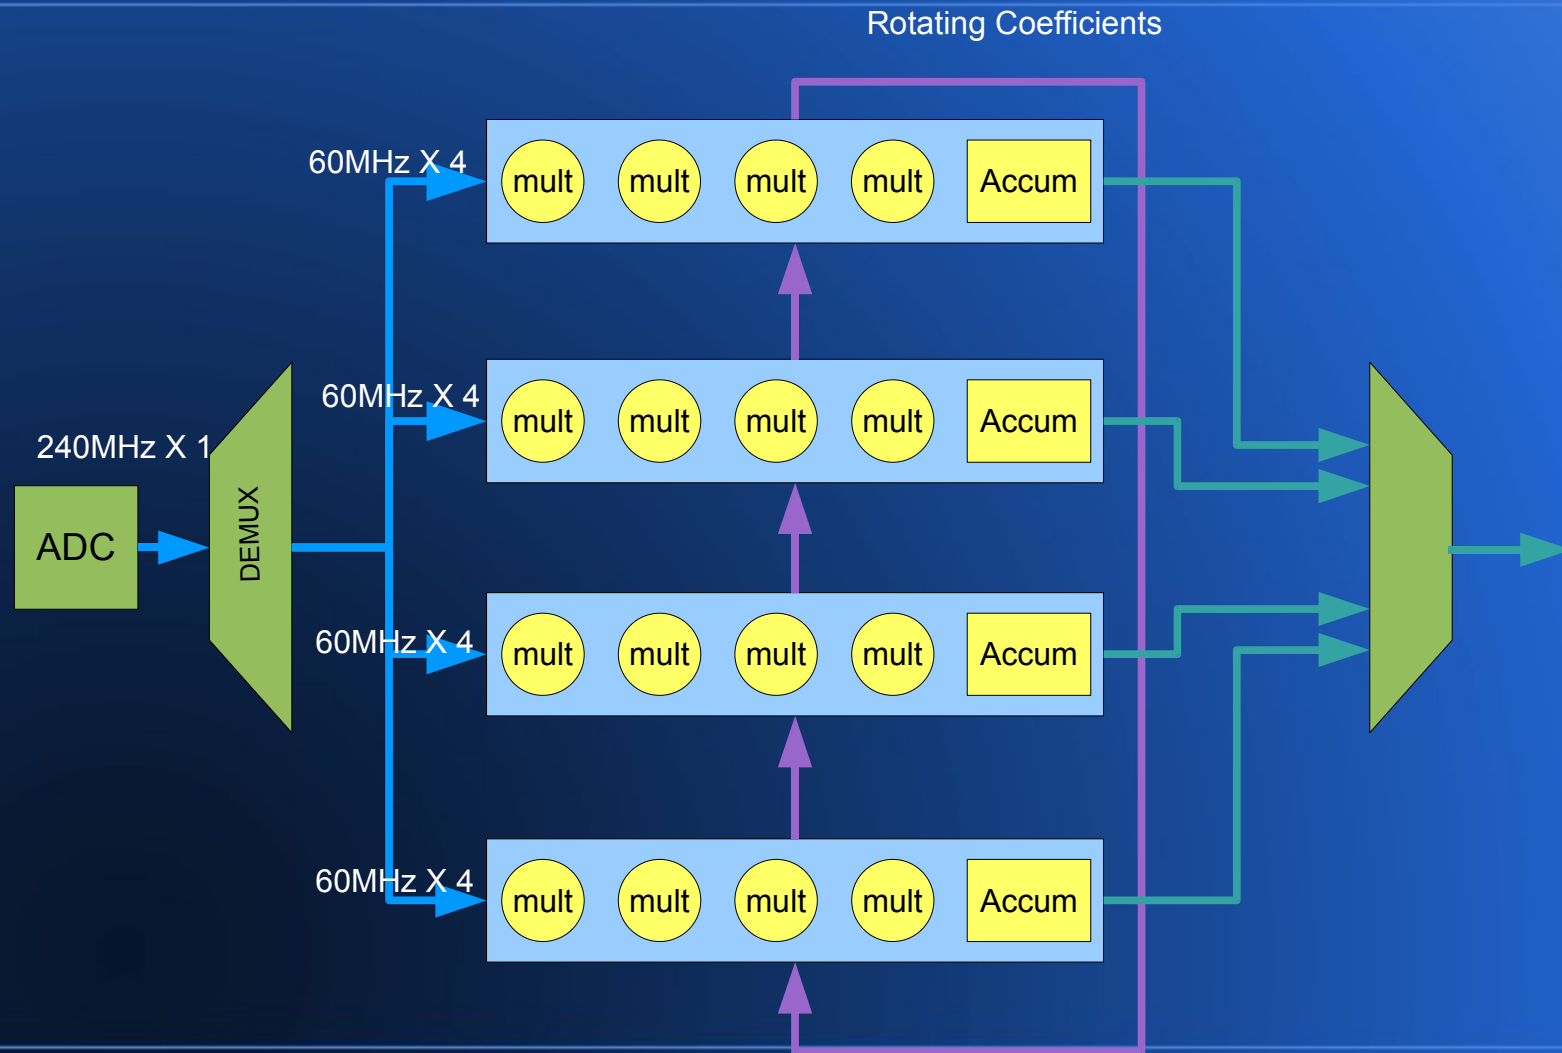
Computed by MAC 1

Computed by MAC 2

Computed by MAC 3

Computed by MAC 4

Our Approach – Thread Oriented



FIR Comparisons

	Resource Required	Pros	Cons
Plain FIR + Decimation	(N/R) Mult.	Straightforward Implementation	Wasted calculations due to decimation
Simple Polyphase FIR	(N/M) Mult.	No wasted calculations	Cannot take advantage of high R
Serial FIR	2^K Entry LUT	Small, compact FIR	Slow, need buffer for high input rates
ISAAC FIR	$(N/(M \cdot R))$ Mult.	Easy implementation and analysis, low resource requirement	Uses more multipliers than Serial FIR

Coefficient Bank Implementation

- Option 1 – Shift Registers
 - XST can map automatically V2 SRL16 macros, save flip-flops for pipeline stages
 - Tuning is done by shifting 1's and 0's in proper sequence
 - However SRL16 cannot be scrubbed
 - Mitigate SEU with EDAC or TMR only

Coefficient Bank Implementation

- Option 2 – BlockRAM
 - Coefficient banks usually small, wasteful to use entire BlockRAM
 - Tuning is done by writing to the BlockRAM
 - Scrubbing can be done by enabling BlockRAM scrubbing
 - However unclear how to selectively scrub BlockRAM, i.e. other BlockRAM used as buffers should not be scrubbed

Coefficient Bank Implementation

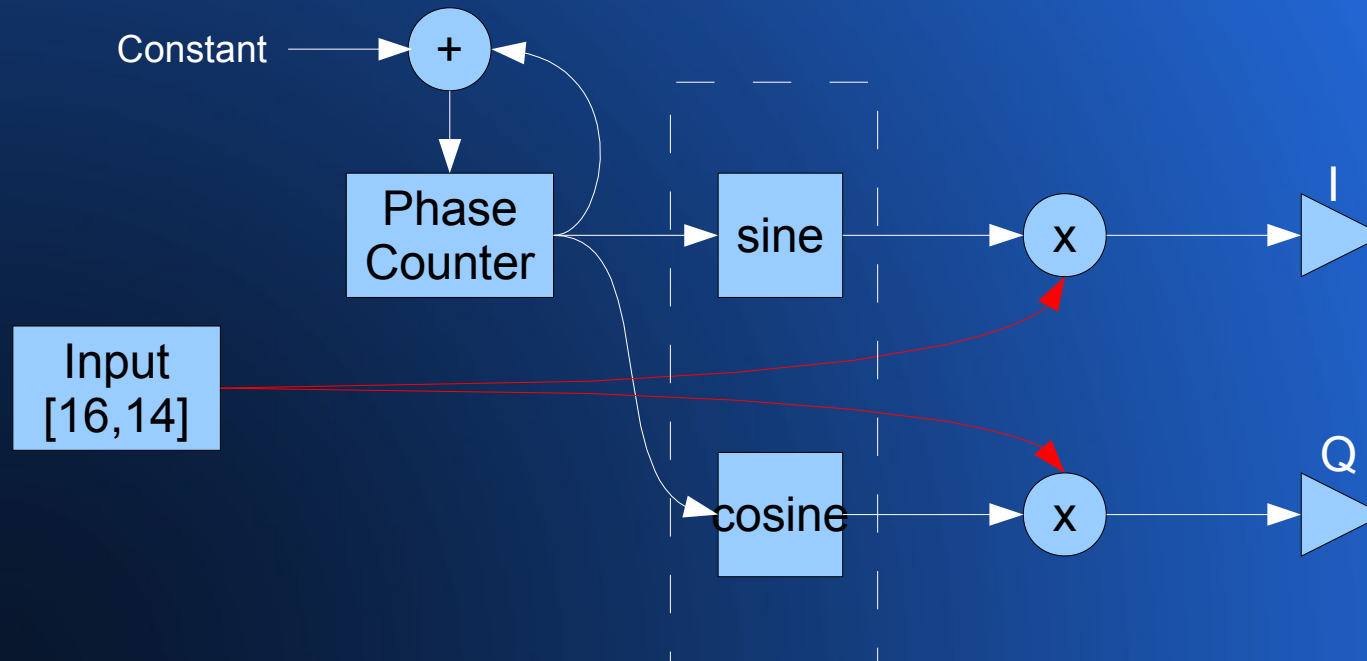
- Option 3 – Distributed RAM (LUTRAM)
 - Coefficient banks usually small enough to fit in distributed RAM
 - Tuning is done by writing to the RAM block
 - Scrubbing is free: distributed RAM is scrubbed together with config RAM.
 - Similar LUT usage as Option 1

Coefficient Bank Summary

	Resource Element	SEU Mitigation	Tuning
Shift Registers	SRL16 (4LUTs)	Cannot scrub; EDAC or TMR only	Bit-wise shift in new values
BlockRAM	BRAM Blocks	BRAM scrubbing or EDAC	Rewrite BRAM content
Distributed RAM	4LUTs	Config. RAM scrubbing or EDAC	Rewrite RAM content

QM Implementation

```
phase=wbpc*(0:N-1)=2pi*fc/f0*(0:N-1)
xbsr=xov.*cos(phase)
xbsi=xov.*sin(phase)
```



QM Implementation

- Phase counter is synchronized/reset on TX_START
- Phase counter should not accumulate quantization error
- Actual implementation uses the following equation (x,y are outputs from 2 FIR0):
 - $I = x \cdot \sin(k) + y \cdot \cos(k)$
 - $Q = x \cdot \cos(k) + y \cdot \sin(k)$
- How to implement the sine/cosine elements?

Sine/Cosine Implementation

- Option 1 – CORDIC
 - Multi-cycle iterative evaluation process
 - Need multiple clock cycles to get the result
 - Latency proportional to precision requirement
 - Two sources of error:
 - phase input quantization error (caused by input quantization)
 - output quantization error (caused by output quantization)

Sine/Cosine Implementation

- Option 2 – Full Sine/Cosine Table
 - Equivalent to CORDIC but implemented as large look-up table
 - Size is 2^K
 - Latency is one clock
 - Two sources of error: phase input quantization error and output quantization error
 - Most of the table is not used!

Sine/Cosine Implementation

- Option 3 – Condensed Sine/Cosine Table
 - $\text{phase} = \text{wbpc} * (0:N-1) = 2\pi * f_c / f_0 * (0:N-1)$
 - Observe that if f_c is constrained to multiples of 1MHz, phase can only take on finite set of values
 - Size is shrank from (2^K) to $(f_0 / 1\text{MHz})$
 - Latency is still one clock
 - Quantization error only on output; phase input quantization error is eliminated!

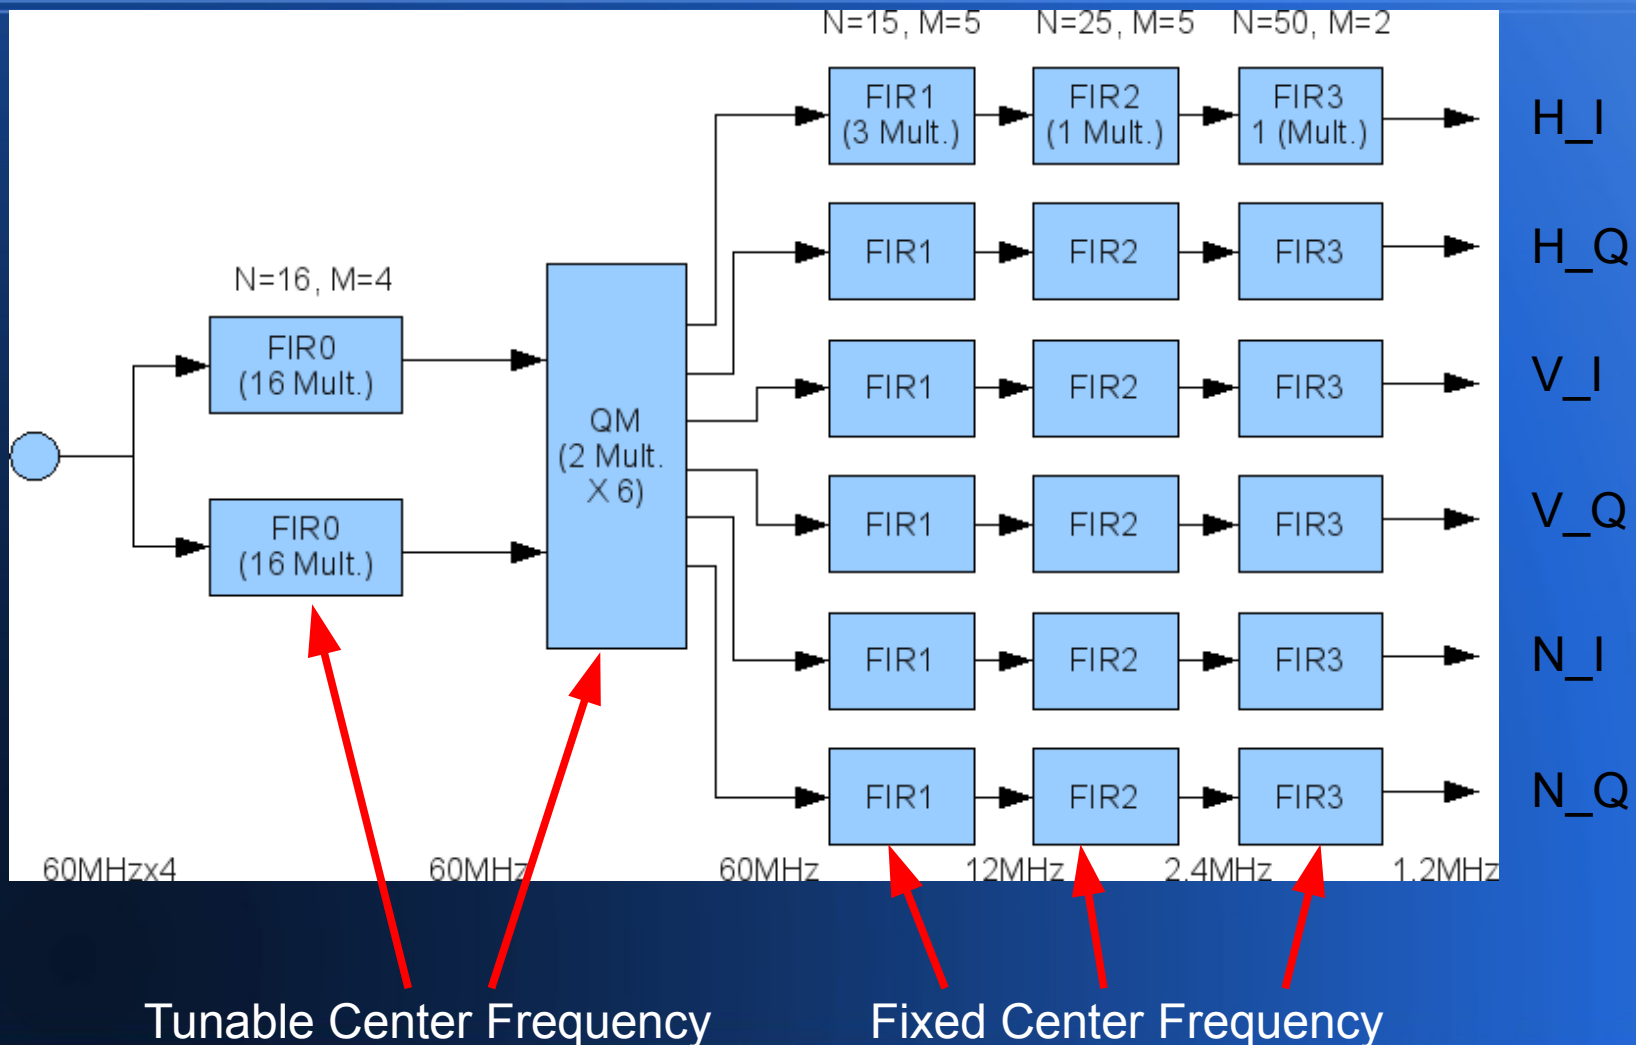
Sine/Cosine Implementation

- Consideration for SEU Mitigation
 - Tables can be implemented with BlockRAM or Distributed RAM
 - Turning on scrubbing for BlockRAM may not be desirable if other modules use BlockRAM as buffers
 - Distributed RAM is scrubbed along with other 4LUT configuration RAM
 - EDAC is also possible, but more overhead and latency

Sine/Cosine Summary

	Size	Latency	Q. Error	SEU Mitigation
CORDIC	small	20+ clock cycles	Phase Input + Function Output	Config RAM scrubbing
Full Table	2 ^k entries (1024)	1 clock cycle	Phase Input + Function Output	BRAM scrubbing
Condensed Table	f ₀ / 1MHz (240)	1 clock cycle	Function Output Only	BRAM or Distributed RAM Scrubbing


FIR and QM Tuning



Tuning FIR0

- Need to load the coefficient banks (32 loads)
- Need to re-compute the coefficients
 - $wbpc = 2 \cdot \pi \cdot f_{chanc} / fs_0$;
 - $hmsbpr = hmsbp \cdot \cos(wbpc \cdot [0:1:hlnghbp-1])$;
 - $hmsbpi = hmsbp \cdot \sin(wbpc \cdot [0:1:hlnghbp-1])$;
- Sine and cosine functions are implemented using the same LUT used in QM, assuming f_{chanc} is multiples of 1MHz

Explanation of f_chanc in SMAP context



- Fcv and Fch are 5MHz apart
- Fcn is 2.5MHz away from both Fcv and Fch
- F_chanc for FIR0 is Fcn
- 1MHz constraint is only on Fcn, not on Fcv and Fch

Tuning QM

- Recall QM phase calculation:
 - $\text{phase} = \text{wbpc} * (0:N-1) = 2\pi * f_c / f_0 * (0:N-1)$
- Tuning QM's f_c only involves changing the constant increment value for the phase counter
- If $f_c = 1\text{MHz}$, increment value is 1
- If $f_c = 34\text{MHz}$, increment value is 34

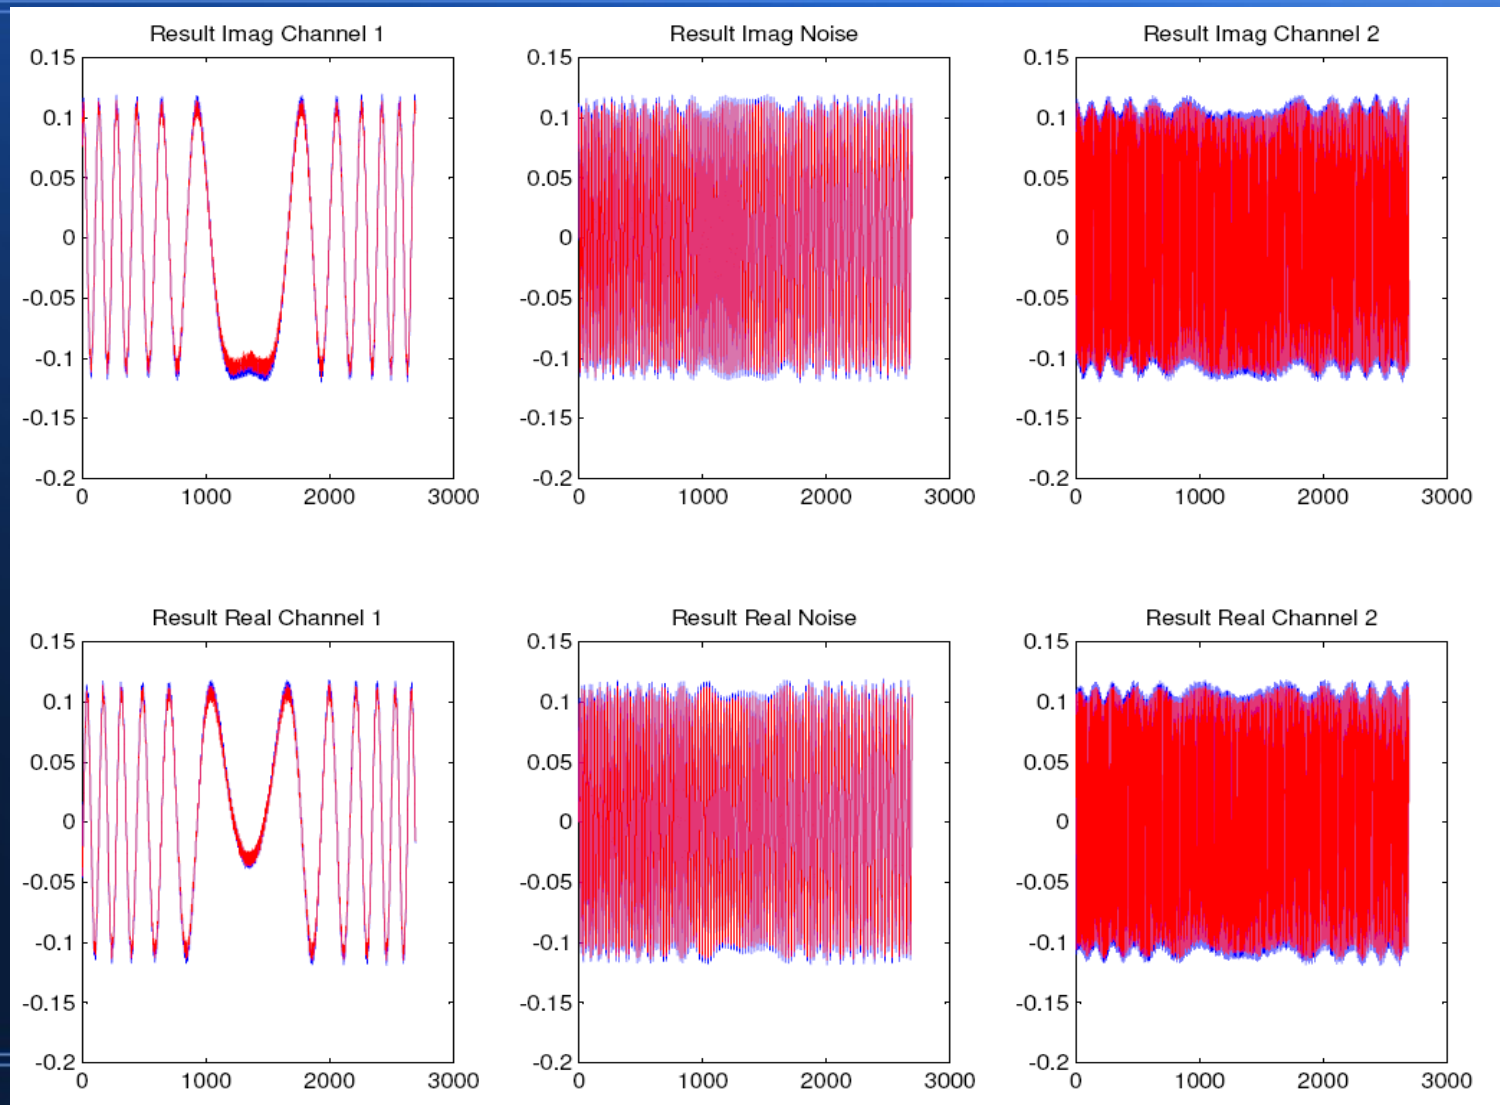
Design Verification

Output Comparison of FIR0 and QM

Output after
FIR0 and
QM
processing

Red – RTL
simulation

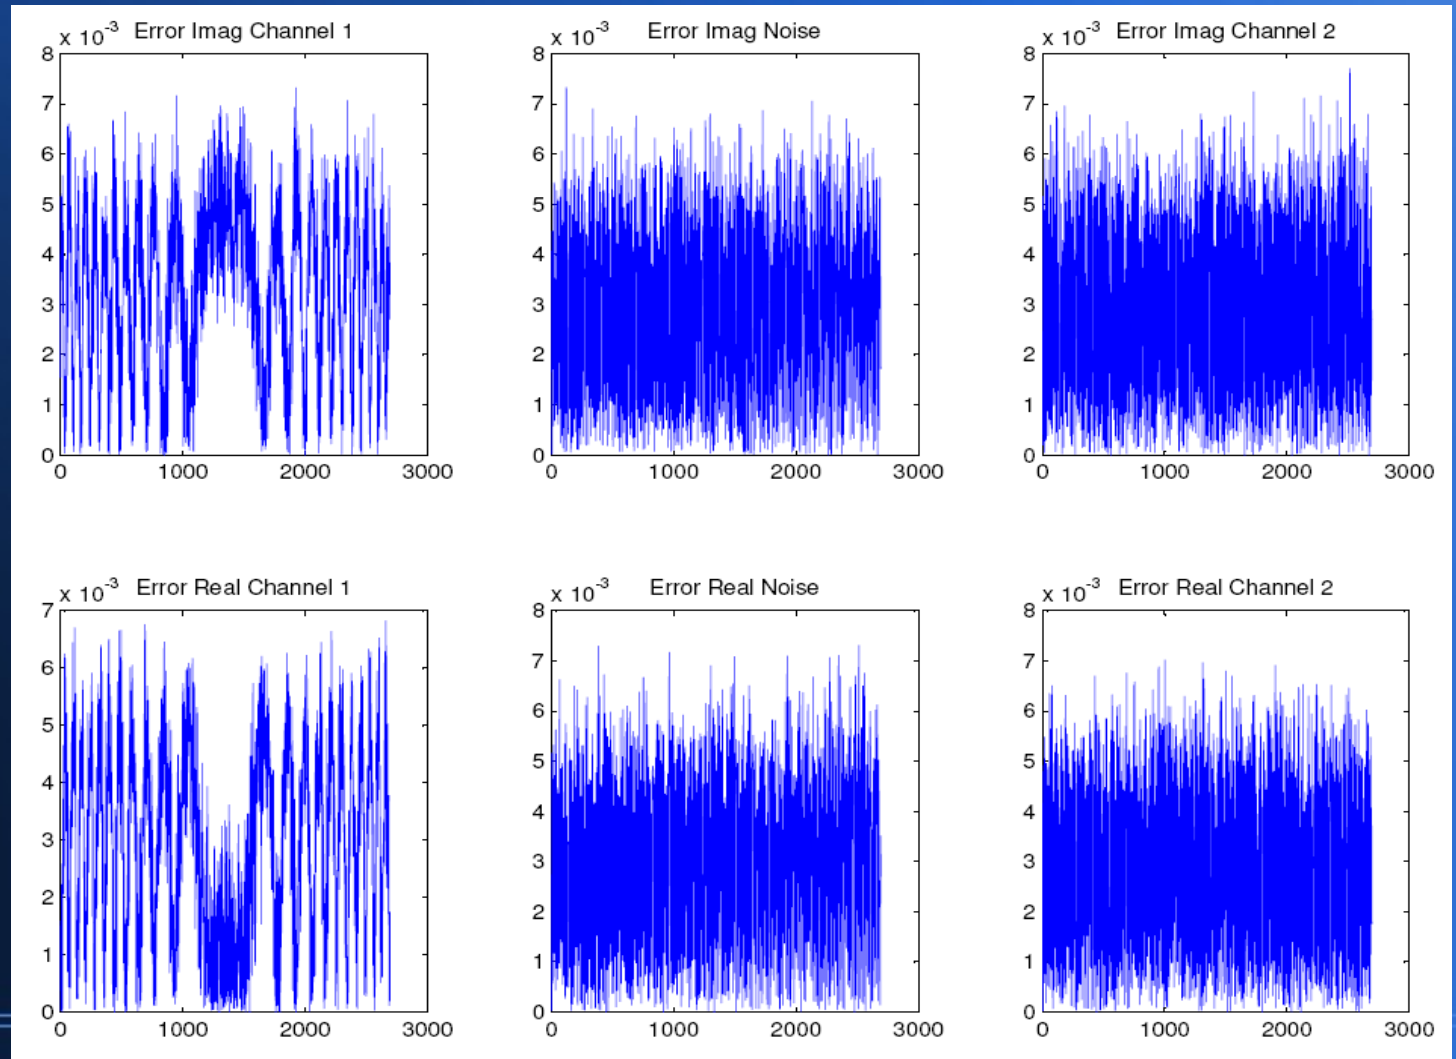
Blue –
Matlab
simulation



Design Verification

Error of FIR0 and QM


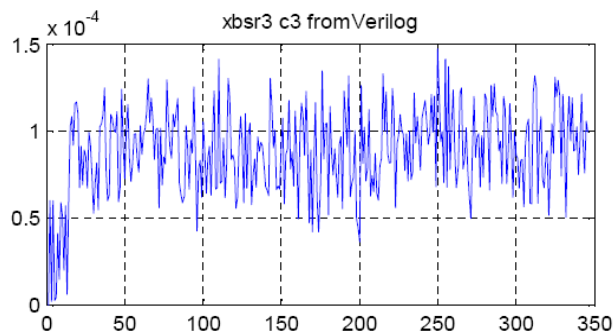
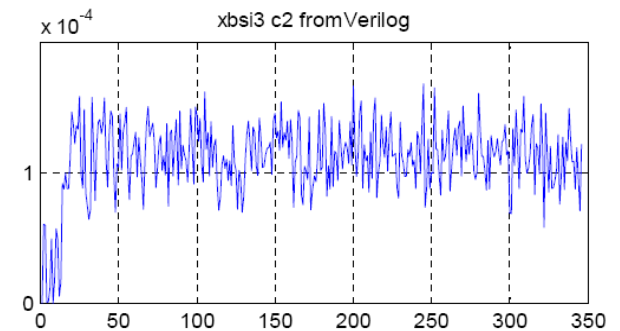



Error after
FIR0 and
QM
processing



Design Verification

Error of FIR1, FIR2, FIR3

Error after
FIR1,
FIR2, and
FIR3
processing



Synthesis Result

Design	240MHz Design	
FPGA Device	Xilinx Virtex 2: speed grade 4 (space grade)	
Speed (Bandwidth)	84.9MHz	
	Used / Total	Percentage %
Number of Multipliers	76 / 96	79%
Number of Slices	8,247 / 14,336	57%
Number of Flip Flops	9,736 / 28,672	33%
Number of 4 input LUTs	13,868 / 28,672	48%

- Resource usage includes FIR0, QM, FIR1, FIR2, and FIR3 for complex data processing for H-pol, V-pol and noise

Future Work

- Further reduce multiplier count with VKCMs (constant multipliers implemented in LUTs)
- FIR0 can save 16x2 full multipliers by replacing with VKCMs
- How to detect error in computation?
- How to selectively scrub BRAMs?